

What's the Main Goal Here?

We're building a **machine learning model** that can “see” an image of a handwritten digit (0–9) and **predict which digit it is**.

So basically, we’re teaching the computer to recognize numbers — the same way your brain recognizes them instantly.

That’s what’s magical here  — making computers see like humans.

The Big Idea: CNN (Convolutional Neural Network)

A CNN is a special kind of neural network built for **images**.

While a regular neural network works on *flat vectors* (1D data), CNNs are designed for **2D structured data** — like an image with pixels arranged in height × width format.

Here’s what CNNs do conceptually:

1. **Convolutional Layers** — Detect *features* like edges, corners, and patterns.
2. **Pooling Layers** — Simplify the image (reduce size) but keep important features.
3. **Fully Connected Layers** — Use those features to classify or predict labels.

Think of CNNs like:

Eyes  (Convolution) → Focus  (Pooling) → Decision  (Fully Connected)

Step-by-Step Concept Breakdown

Dataset — MNIST

The **MNIST dataset** is a collection of 70,000 handwritten digits:

- 60,000 for training
- 10,000 for testing

Each image is **28×28 pixels**, grayscale (1 channel).

We’re training our CNN to map:

Image pixels → Correct digit (0–9)

2 Data Loading

We use:

```
torchvision.datasets.MNIST(...)
```

to automatically download and load the data.

Each image is converted into a **tensor** (PyTorch format for numerical data).

`transforms.ToTensor()` converts each pixel from [0,255] to [0,1], making it easier for the model to learn.

`DataLoader` helps to:

- Split data into **batches** (64 images at a time)
 - Shuffle data each epoch (prevents overfitting)
-

3 CNN Architecture — The Brain of It All

Let's break your CNN class conceptually:

Conv2d(1, 16, kernel_size=3, padding=1)

- Input: 1 channel (grayscale)
- Output: 16 feature maps
- Kernel (3×3): slides over image to detect small patterns (edges, curves)
- Padding keeps image size same.

This layer **learns to detect local features** — things like strokes or small lines.

MaxPool2d(2, 2)

- Reduces each feature map's size by 2×.
- Keeps only *most important features* (maximum value in a region).

Imagine zooming out but keeping key patterns visible.

Conv2d(16, 32, kernel_size=3, padding=1)

- Takes the first 16 learned features and combines them to detect *more complex* features (like loops, edges, corners).

So by the second convolution, the model starts recognizing digit-specific shapes.

Linear(32 * 7 * 7, 128) & Linear(128, 10)

After flattening, these are **fully connected layers**.

Here the network connects all learned features to output probabilities for **10 digits**.

Forward Propagation — How Data Flows

```
x = self.pool(F.relu(self.conv1(x)))  
  
x = self.pool(F.relu(self.conv2(x)))  
  
x = x.view(-1, 32 * 7 * 7)  
  
x = F.relu(self.fc1(x))  
  
x = self.fc2(x)
```

Each layer transforms the data step-by-step:

1. Convolution detects patterns.
2. ReLU (activation) adds non-linearity — helps model learn complex relationships.
3. Pooling compresses data.
4. Fully connected layers combine all features to classify.

At the end, you get 10 values → representing probabilities of each digit (0–9).

Training — How the Model Learns

We're doing *supervised learning* here.

That means we already know the correct digit (label) for each image.

The process:

1. Feed image into CNN → get prediction.
2. Compare prediction with true label using **loss function**.
3. Adjust weights using **backpropagation** to reduce loss.

Components:

- **Loss function:** CrossEntropyLoss() — measures how wrong the model's predictions are.
 - **Optimizer:** Adam — adjusts weights intelligently using gradient descent.
 - **Epochs:** Number of complete passes through training data.
-

Evaluation — Testing the Model

After training, we evaluate on test data it hasn't seen before.

We calculate:

$\text{correct} / \text{total} * 100$

to get **accuracy** — how often our model correctly predicts the digit.

This shows whether the CNN has **actually learned to generalize**, not just memorize.

The Core Concepts Behind It

Concept	Meaning
Convolution	Extracts features from the image using small filters
Activation (ReLU)	Adds non-linearity, helping learn complex patterns
Pooling	Reduces spatial size and helps prevent overfitting
Flattening	Converts 2D feature maps to 1D vector for fully connected layers
Fully Connected Layers	Combine features to make predictions
Backpropagation	Updates weights using gradients
Cross-Entropy Loss	Measures difference between predicted and true labels

What We're Proving / Demonstrating Here

This whole experiment **proves that a CNN can automatically learn visual patterns from raw pixels** — no manual feature engineering required.

We're showing that:

- Machines can see and *classify* handwritten digits accurately.
- CNNs outperform traditional models because they understand **spatial structure** of images.
- Deep learning truly *learns hierarchical representations* (edges → shapes → digits).